

Entity Component Systems for Compilers

Reimagining the Abstract Syntax Tree

or

“How not to build a compiler”

J Pratt
they/them

Arcs (Sydney)
Raksha

[jopra@
Cypher1](mailto:jopra@Cypher1)

What on earth is an ECS?

Entity Component System

Used in Game development

Alternative to Object Oriented Programming

What do I mean by ECS?

Entity

A unique object (e.g. a character)

What do I mean by ECS?

Entity **Component**

A group of properties / values
(e.g. A character's items, skills, etc.)

What do I mean by ECS?

Entity Component **System**

A structured way to handle the components

e.g.

A renderer,

A physics simulator,

A damage system

What do I mean by ECS?

Let's compare to OOP

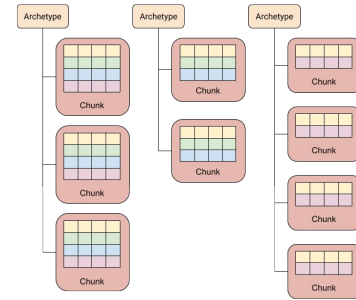
What do I mean by ECS?

Extra terminology:

Archetype = A group of components

System = A thing that processes entities

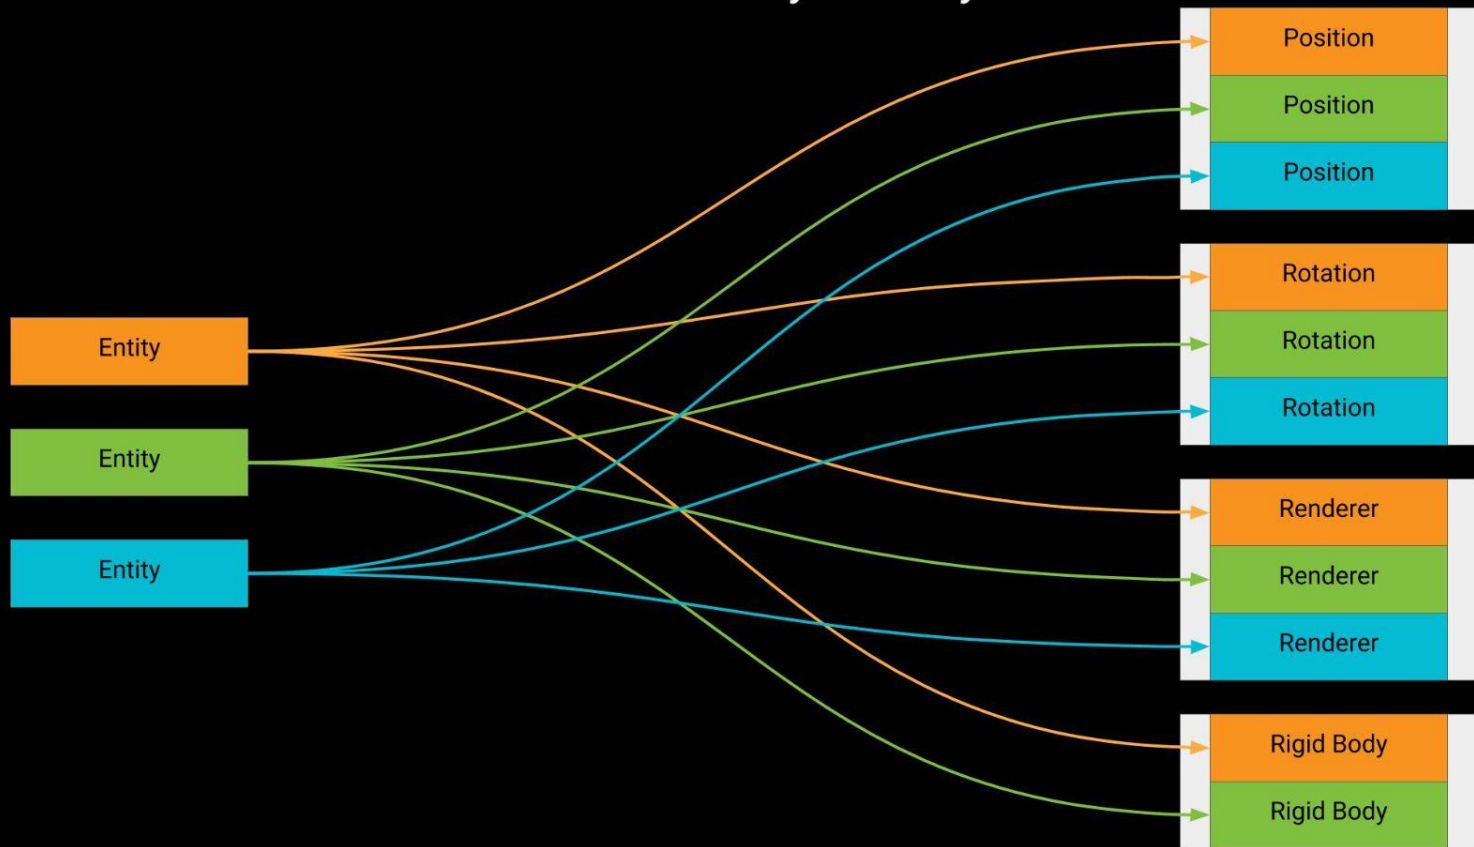
Chunk = Group of co-located instances of the same component/archetype



OOP → A GameObject is a container.



ECS → An entity is a key.



What is an ECS like?

1. Systems operate over components
2. Components are co-located in memory
3. Look up time for each component is small (normally $O(1)$)
4. Systems perform the same (or similar) actions on each component

Why is ECS used in Game development?

1. Separates concerns
2. Systems can operate independently
 - a. Systems can operate in parallel*
 - b. Systems can run at different times
3. Cache & Memory locality
4. Avoiding branches

Sources:

<https://rams3s.github.io/blog/2019-01-09-ecs-deep-dive/> <- This one is probably the best imo

<https://www.richardlord.net/blog/ecs/why-use-an-entity-framework.html>

<https://www.gamedevs.org/uploads/data-driven-game-object-system.pdf>

<https://web.archive.org/web/20171030021158/http://entity-systems-wiki.t-machine.org/>

<http://gameprogrammingpatterns.com/component.html>

*separable systems: i.e. those using non-overlapping sets of components and global data


How is program data stored in compilers?

Abstract Syntax Tree!

AST for short

What do I mean by AST?


Abstract Syntax Tree



What do I mean by AST?

Abstract vs Concrete


We're happy to throw away some
information



What do I mean by AST?

Abstract **Syntax** vs Semantics


This doesn't (normally) model the
runtime behaviour



What do I mean by AST?

Abstract Syntax **Tree** vsGraph?


They (typically) only branch out...





How is program data stored in compilers?

e.g.

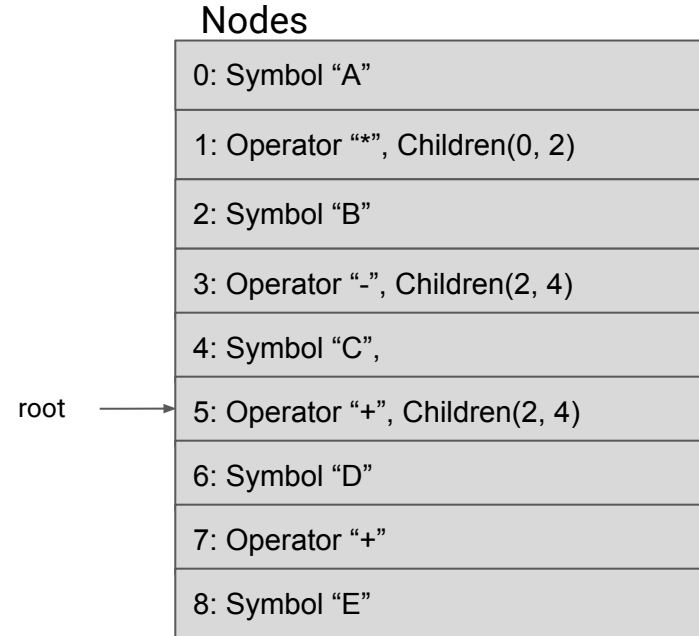


$A * (B - C) + D + E$




How is program data stored in compilers?



How is program data stored in compilers?




What would this look like in an ECS?



What would this look like in an ECS?

IsSymbol


A
B
C
D
E




What would this look like in an ECS?

IsSymbol

0: A
1: B
2: C
3: D
4: E



What would this look like in an ECS?




IsSymbol

0: A
1: B
2: C
3: D
4: E

Entities

0: IsSymbol 0
2: IsSymbol 1
4: IsSymbol 2
6: IsSymbol 3
8: IsSymbol 4

What would this look like in an ECS?




IsSymbol

0: E0, A
1: E2, B
2: E4, C
3: E6, D
4: E8, E

Entities

0: IsSymbol 0
2: IsSymbol 1
4: IsSymbol 2
6: IsSymbol 3
8: IsSymbol 4

What would this look like in an ECS?




IsSymbol

0: E0, A
1: E2, B
2: E4, C
3: E6, D
4: E8, E

Entities

0: IsSymbol 0
1: IsOperator 0
2: IsSymbol 1
3: IsOperator 1
4: IsSymbol 2
5: IsOperator 2, IsRoot
6: IsSymbol 3
7: IsOperator 3
8: IsSymbol 4

What would this look like in an ECS?



IsOperator

*
-
+
+

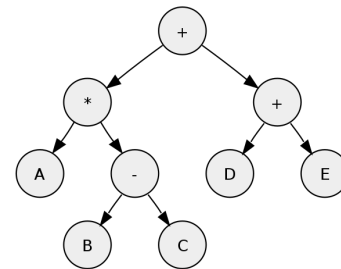
IsSymbol

0: E0, A
1: E2, B
2: E4, C
3: E6, D
4: E8, E

Entities

0: IsSymbol 0
1: IsOperator 0
2: IsSymbol 1
3: IsOperator 1
4: IsSymbol 2
5: IsOperator 2, IsRoot
6: IsSymbol 3
7: IsOperator 3
8: IsSymbol 4

What would this look like in an ECS?



IsOperator

0: *
1: -
2: +
3: +


IsSymbol

0: E0, A
1: E2, B
2: E4, C
3: E6, D
4: E8, E

Entities

0: IsSymbol 0
1: IsOperator 0
2: IsSymbol 1
3: IsOperator 1
4: IsSymbol 2
5: IsOperator 2, IsRoot
6: IsSymbol 3
7: IsOperator 3
8: IsSymbol 4

What would this look like in an ECS?



IsOperator

0: E1, *
1: E3, -
2: E5, +
3: E7, +


IsSymbol

0: E0, A
1: E2, B
2: E4, C
3: E6, D
4: E8, E

Entities

0: IsSymbol 0
1: IsOperator 0
2: IsSymbol 1
3: IsOperator 1
4: IsSymbol 2
5: IsOperator 2, IsRoot
6: IsSymbol 3
7: IsOperator 3
8: IsSymbol 4

What would this look like in an ECS?



IsOperator

0: E1, *, E0, E3
1: E3, -, E2, E4
2: E5, +, E1, E7
3: E7, +, E6, E8


IsSymbol

0: E0, A
1: E2, B
2: E4, C
3: E6, D
4: E8, E

Entities

0: IsSymbol 0
1: IsOperator 0
2: IsSymbol 1
3: IsOperator 1
4: IsSymbol 2
5: IsOperator 2, IsRoot
6: IsSymbol 3
7: IsOperator 3
8: IsSymbol 4

What would this look like in an ECS?



IsOperator

0:	E1, *, E0, E3
1:	E3, -, E2, E4
2:	E5, +, E1, E7
3:	E7, +, E6, E8

IsSymbol

0:	E0, A
1:	E2, B
2:	E4, C
3:	E6, D
4:	E8, E


IsValue

....
: E, 123
...

Entities

0:	IsSymbol 0
1:	IsOperator 0
2:	IsSymbol 1
3:	IsOperator 1
4:	IsSymbol 2
5:	IsOperator 2, IsRoot
6:	IsSymbol 3
7:	IsOperator 3
8:	IsSymbol 4

What would this look like in an ECS?



IsOperator

0: E1, *, E0, E3
1: E3, -, E2, E4
2: E5, +, E1, E7
3: E7, +, E6, E8

IsSymbol

0: E0, A
1: E2, B
2: E4, C
3: E6, D
4: E8, E

IsValue

....
: E, 123
...

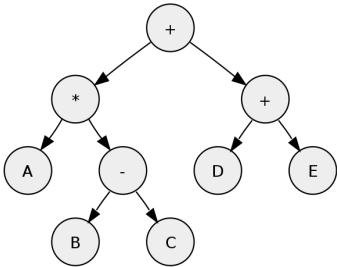
SourceLocation

....
: E, "test.rs" L5:C3
...

Entities

0: IsSymbol 0
1: IsOperator 0
2: IsSymbol 1
3: IsOperator 1
4: IsSymbol 2
5: IsOperator 2, IsRoot
6: IsSymbol 3
7: IsOperator 3
8: IsSymbol 4

How is program data stored in compilers?



AST

Nodes

0: Symbol "A"
1: Operator "*", Children(0, 2)
2: Symbol "B"
3: Operator "-", Children(2, 4)
4: Symbol "C",
5: Operator "+", Children(2, 4)
6: Symbol "D"
7: Operator "+"
8: Symbol "E"

root →

ECS

IsOperator

0: E1, *, E0, E3
1: E3, -, E2, E4
2: E5, +, E1, E7
3: E7, +, E6, E8

IsSymbol

0: E0, A
1: E2, B
2: E4, C
3: E6, D
4: E8, E

IsValue

....
: E, 123
...

SourceLocation

....
: E, "test.rs" L5:C3
...

Entities

0: IsSymbol 0
1: IsOperator 0
2: IsSymbol 1
3: IsOperator 1
4: IsSymbol 2
5: IsOperator 2, IsRoot
6: IsSymbol 3
7: IsOperator 3
8: IsSymbol 4

Why might we want ECS in Compilers?

1. Separates concerns
2. Systems can operate independently
 - a. Systems can operate in parallel*
 - b. Systems can run at different times
3. Cache & Memory locality
4. Avoiding branches

What can ECS do that might work in Compilers?

Anything that can be written as repeated passes over related nodes

- This seems plausible for Static analysis!

So:

- Optimization!
- Type checking!
- Dataflow checking!
- Code generation!

A lot of what a modern compiler spends time on!

Does it actually work?

Well...

I don't know.

Let's write some benchmarks!

Steel: An attempt at ECS for compilers

- A small expression language
- Just enough program to do optimization on
- ...
- Okay, maybe also a few fun features... just to keep life interesting

Steel: An attempt at ECS for compilers

Features

- Parser
 - Precedence:
i.e. $1+2*3 = 2*3+1 = 6$
 - Keyword arguments: `f(x=3)`
 - Positional arguments: `f(3)`
- Pretty printer
- Optimizer
 - Does constant folding
e.g. $1+2 \Rightarrow 3$
- Interpreter
 - So many (two) types!
 - `BuiltInFunction`
 - `i64`
- and a random code generator

Steel: An attempt at ECS for compilers

Features

- Implemented in AST **AND** ECS
- 100%* Rust 🦀
- **Has Benchmarks**

*(ignore the data collection scripts)

Steel: An attempt at ECS for compilers

Demo?

Steel: An attempt at ECS for compilers

Benchmarks

- Timed operations over generated programs of arbitrary size
- Used two different storage backends with an abstraction layer
 - AST
 - ECS
- Programs were ‘trivially optimizable’

E.g.


$1+2+3+4+\dots+1000$

$1*2*3*4+\dots*1000$

Steel: An attempt at ECS for compilers

Benchmark results

(shorter bar = faster)

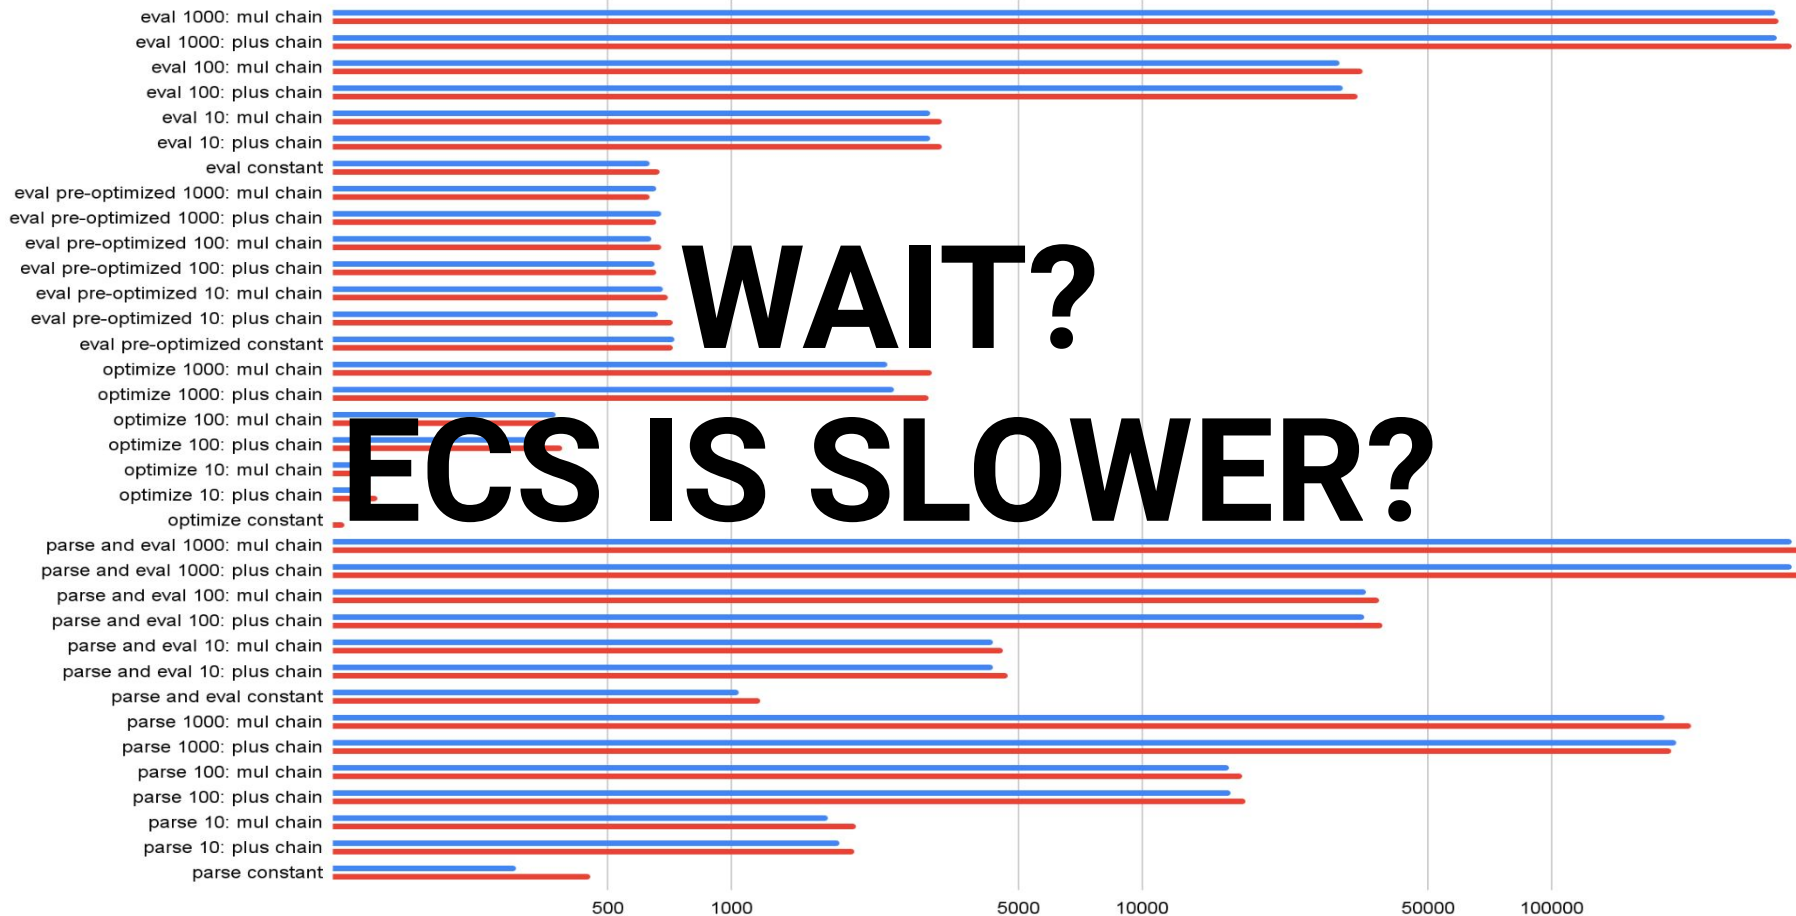


See the appendices
for an update

■ ast avg ■ ecs avg



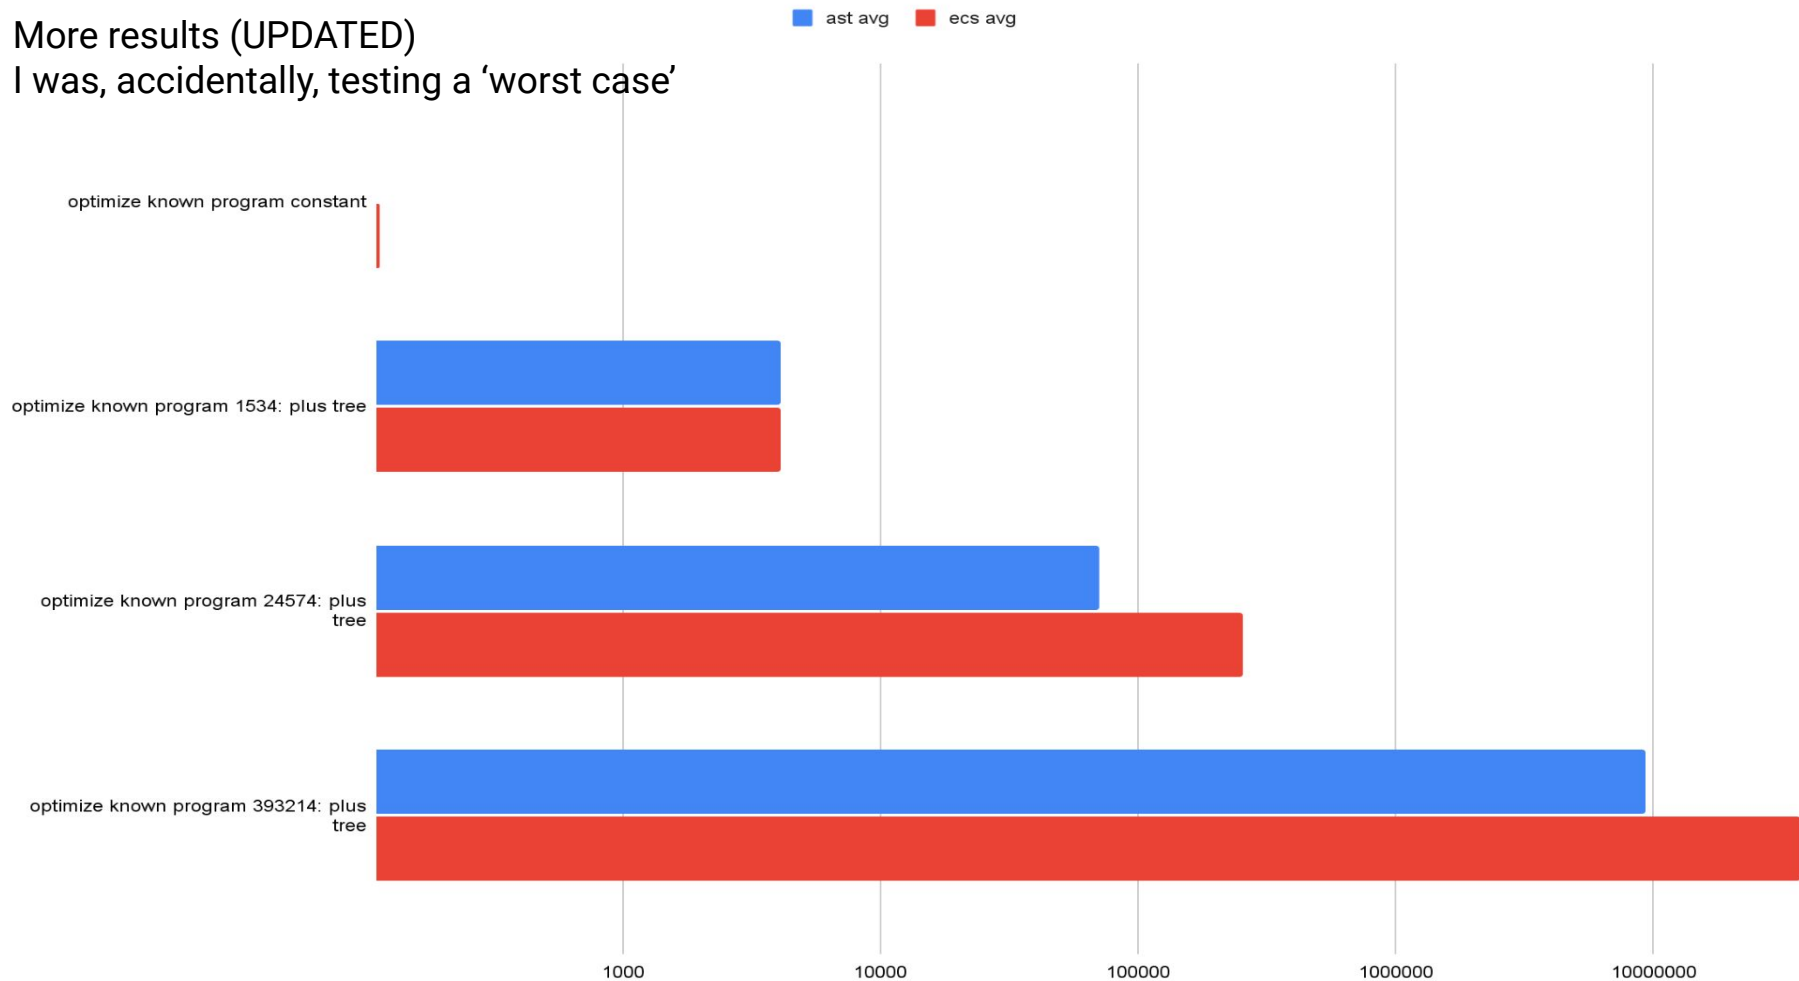
ast avg ecs avg



(shorter bar = faster)


More results (UPDATED)

I was, accidentally, testing a 'worst case'



More results (UPDATED)

ast avg ecs avg



Steel: What could be happening here?

‘Sins’ I committed when preparing this data

I.e. Places where mistakes may have been made.

- Used `bash` to process the output of Rust’s benchmark data
 - because it there was so much
 - because manually converting the data into a sheet/excel was way too much work
- Used sheets to convert units with a handwritten function
 - because sheets doesn’t understand small units of time

Steel: What could be happening here?

- My code could be really really bad

Steel: What could be happening here?

- My code **IS** suboptimal
- My benchmarks could be bad
- My ECS implementation **IS** slow
- My AST implementation... could be **fast???**

Future work

Future work: Move on?

Maybe AST is good actually?

Future work: Fix it?

- Try other ECS implementations
 - Bevy
 - Specs
- Try different modeling
- Maybe the way I'm using ECS is poor
 - Hmm...

Future work: Maybe the way I'm using ECS is poor

- The optimizer is a big *“Find and Replace”*
- We have good *“Find and Replace”* algorithms!

Future work: Maybe the way I'm using ECS is poor

We have good “*Find and Replace*” algorithms

- Regex??? (Not great for structured data)
- Techniques from Databases? (i.e. optimizing SQL queries)
- [RETE](#)
 - Uses graphs of candidates
 - Avoids re-doing work
 - Might be a good fit
- RETE is used by a bunch of different (non-compiler) tools already
 - Datalog
 - [Souffle](#)
- Maybe I should try building a compiler with these?
- Others?

Future work: Collaboration!

- [Open source Steel](#)
- Write a blog post?
- Write a paper?

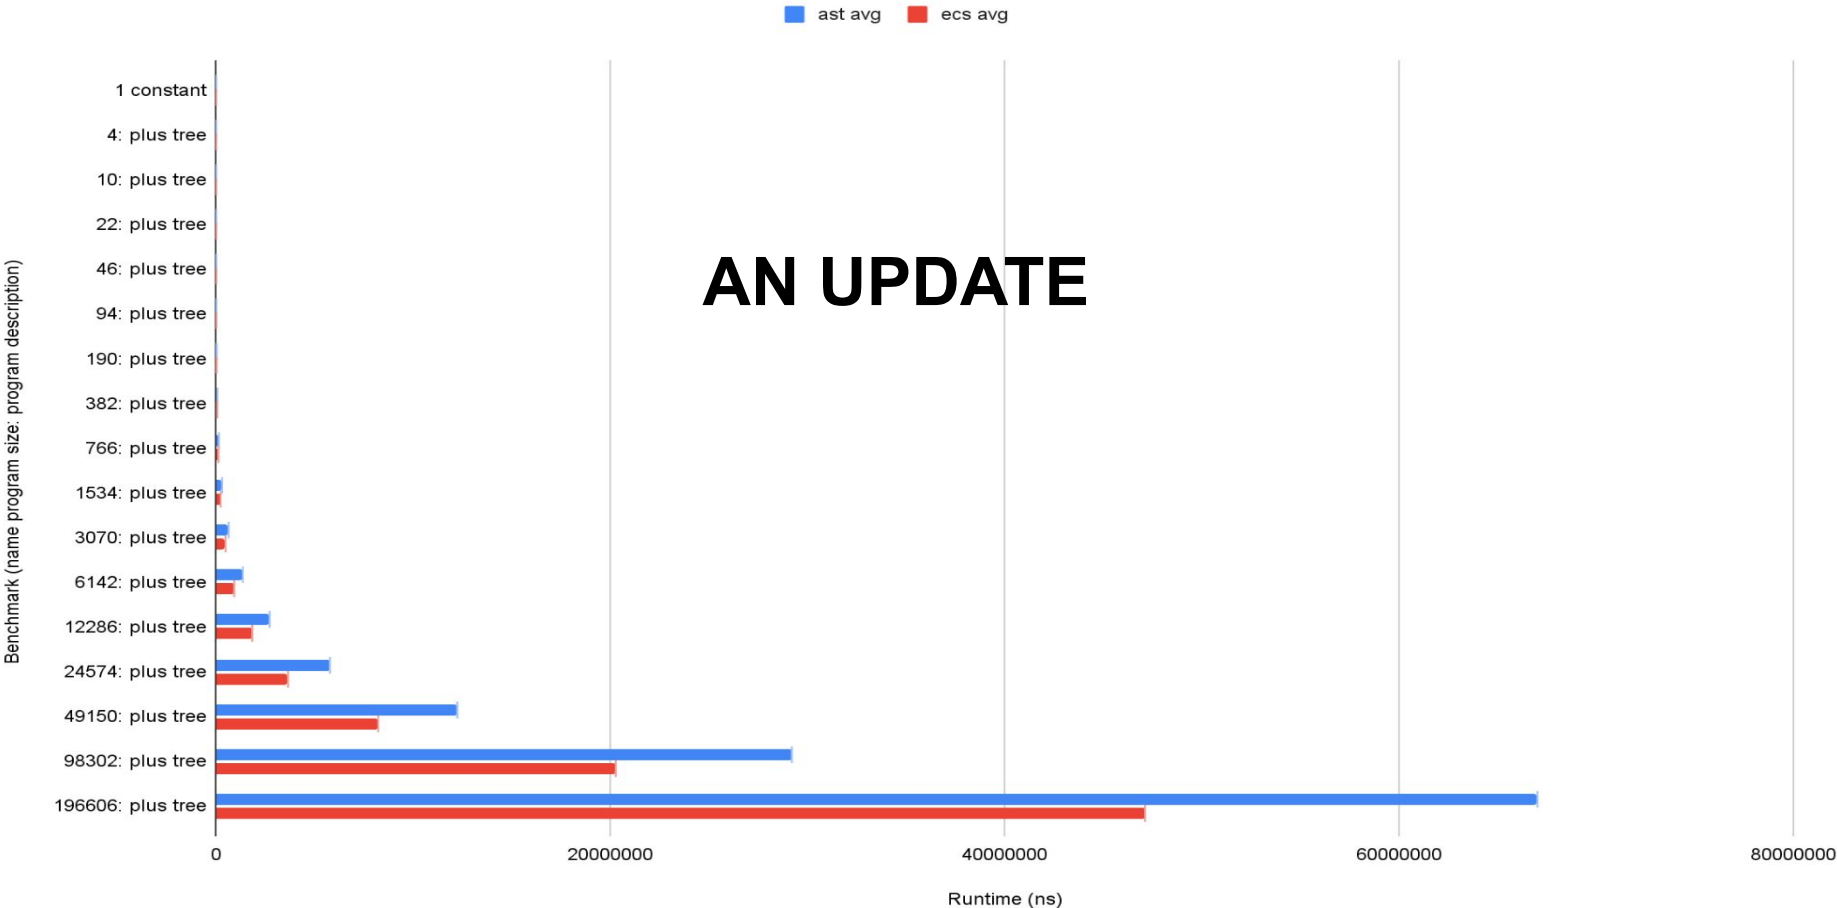
Contributions & Collaborations welcome

:D

Thanks!

Continue on for appendices!

Comparison of runtimes optimizing trivial programs ECS vs AST: note linear scale used



Improvements

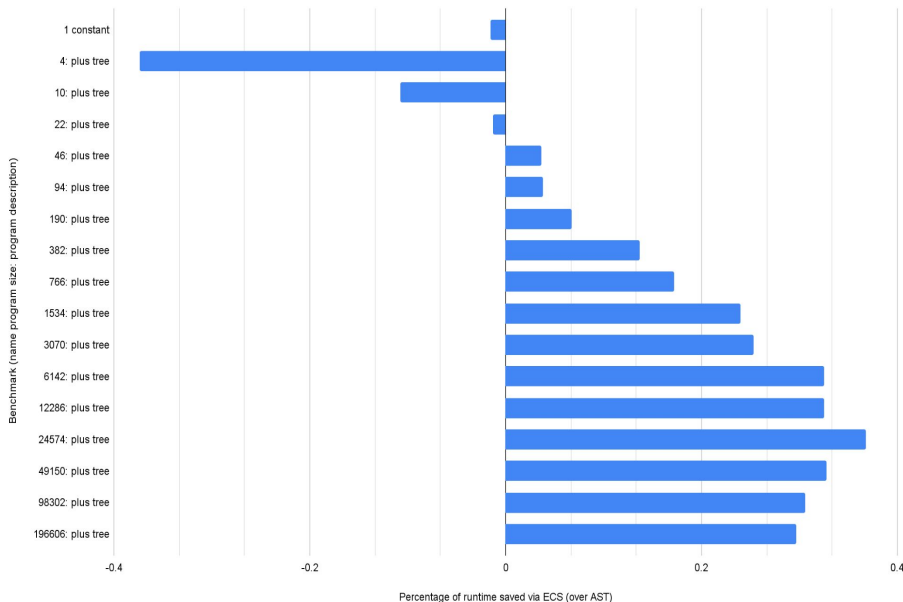
Fixes:

- Split Symbol and Operator tables (like in slides)
- Store 'left' and 'right' with calls
(separate from the names args)
- Fetch Values and Operators from program data,
Don't use a hashmap for bookkeeping

Results:

- On small programs ECS's constant factors outweigh the benefits :(
- On large programs ECS can achieve ~20%-38% performance improvement :)

Time saved by ECS as a percentage of AST run time



Improvements

Takeaway:

- Pre-computation has to be **FASTER** than memory
- HashMaps are awesome, but dense, purpose built data structures are generally better
- Parser is 97% of my benchmark time, even using a library like [nom](#)
 - At this point I **should** probably stop

The following talk was **FAR** more helpful than anything else I watched or read:

- [A Practical Guide to Applying Data-Oriented Design](#) (by Zig creator: Andrew Kelley)

Rust tools I used (or wanted to)

CLI tools

- <https://crates.io/crates/bacon/1.1.4> // A watch command
- [sccache - crates.io: Rust Package Registry](https://crates.io/crates/sccache) // A distributed compilation cache

Useful libraries

- https://crates.io/crates/static_assertions // static checks to avoid making mistakes
- <https://crates.io/crates/memoize> // caching/memoizing
- <https://github.com/salsa-rs/salsa> // incremental computation
- <https://github.com/Geal/nom> // makes parsers easy
- https://docs.rs/typed-arena/latest/typed_arena/ // Makes contiguous memory easier ([comparison](#))

ECS frameworks

- [Bevy](#) // Comes recommended and with a great set of game dev tools
- <https://github.com/amethyst/specs>

Other compiler projects & tools

Tools

- <https://github.com/jzimmerman/langcc> // Tools for generating compilers
- [Rust!](https://www.rust-lang.org/) - <https://www.rust-lang.org/> // My recommendation for a programming language

Compiler Projects

- [Cypher1/tako](#) // My hobby project (very WIP)
- [google-research/raksha](#) // My day job

Approaches and algorithms

Equality Saturation

- [Equality Saturation: a New Approach to Optimization *](#)
- <https://rosstate.org/publications/eqsat/RossThesis.pdf>
- <https://arxiv.org/pdf/2004.03082.pdf>

RETE algorithm

- Wiki: https://en.wikipedia.org/wiki/Rete_algorithm
- [Rete Algorithm - an overview | ScienceDirect Topics](#)
- UNSW: [The RETE Algorithm](#)

Index mistakes are easy to make!

Typed ptrs & indexes can help!

Keeping type information can
avoid mixing up indexes.

Code sample: [here](#)

Gist: [Typed Index and Arena impl](#)

```
use std::marker::PhantomData;

#[derive(Clone, Copy, Eq, PartialEq, Ord)]
struct TypedIndex<T> {
    index: usize,
    ty: PhantomData<T>,
}

impl<T> std::fmt::Debug for TypedIndex<T> {
    fn fmt(&self, f: &mut std::fmt::Formatter) -> std::fmt::Result {
        write!(f, "{}_index({:?})", std::any::type_name::<T>(), self.index)
    }
}

impl<T> PartialEq for TypedIndex<T> {
    fn eq(&self, other: &Self) -> bool {
        self.index == other.index
    }
}

impl<T> TypedIndex<T> {
    fn new(index: usize) -> Self {
        Self { index, ty: PhantomData }
    }
}

fn main() {
    let args: Vec<String> = std::env::args().collect();
    let x = args.len();
    let ind1: TypedIndex<i32> = TypedIndex::new(x);
    let ind2: TypedIndex<i32> = TypedIndex::new(x+1);
    dbg!(std::mem::size_of::<TypedIndex<i32>>() == std::mem::size_of::<usize>());
    dbg!(ind1, ind2);
    dbg!(ind1 == ind2);
}
```

Thanks again!